

This is title

Feng Gu(T00751197), author 2 and author 3

April 15, 2025

**Abstract** Here to attach abstract.

## 1 Bayesian Improved Logistic Regression

In this section, we will try to create a Bayesian improved logistic regression using splines. The idea is to use the splines to model the non-linear relationship between the variables and the response variables. Instead of using the classical statistical methods - Maximum Likelihood Estimation (MLE) to estimate the parameters of the model, we will use the Bayesian approach to estimate the posterior distribution of the parameters. The predicted probability of the target variable being 1 will be a distribution up to the posterior distribution of the parameters and the final predicted label will depend on the expectation of the posterior distribution.

### 1.1 Transformation on design matrix for regression splines

Regression splines is a powerful tool to model the non-linear relationship between the input variables and the response variable. It splits the input space into several intervals and fits basis functions to each interval. The basis functions can be chosen to be simple linear or polynomial functions, or more complex functions. The functions across the intervals are connected at the knots, which makes sure the total regression function is continuous and smooth.

Given a univariate predictor  $x \in R^n$ , we can construct a regression spline design matrix by transforming  $x$  into a set of basis functions. For a spline of degree  $d$  with  $K$  knots  $\{\xi_1, \xi_2, \dots, \xi_K\}$ , the new design matrix  $\mathbf{X}_{\text{spline}}$  is:

$$\mathbb{K} : X(n, p) \mapsto X_{\text{spline}}(n, p + C), \quad C > 1$$

$$\mathbf{X}_{\text{spline}} = \begin{bmatrix} 1 & x_{11} & x_{11}^2 & x_{21} & x_{21}^2 & \cdots & x_{p1}^d & (x_{11} - \xi_{11})_+^d & \cdots & (x_{p1} - \xi_{pK})_+^d \\ 1 & x_{12} & x_{12}^2 & x_{22} & x_{22}^2 & \cdots & x_{p2}^d & (x_{12} - \xi_{11})_+^d & \cdots & (x_{p2} - \xi_{pK})_+^d \\ \vdots & \vdots & \vdots & \vdots & \vdots & & \vdots & \vdots & & \vdots \\ 1 & x_{1n} & x_{1n}^2 & x_{2n} & x_{2n}^2 & \cdots & x_{pn}^d & (x_{1n} - \xi_{11})_+^d & \cdots & (x_{pn} - \xi_{pK})_+^d \end{bmatrix}$$

Here,  $(x - \xi_j)_+^d$  denotes the truncated power basis function:

$$(x - \xi_j)_+^d = \begin{cases} (x - \xi_j)^d & \text{if } x > \xi_j \\ 0 & \text{otherwise} \end{cases}$$

This design matrix allows the regression model to fit a flexible, piecewise polynomial function with continuity at the specified knots.

## 1.2 The likelihood function of $\beta$ in the logistic regression

To a data set having target variable that has only two values, we view the target variable follows a Bernoulli distribution with true parameter  $p$ . The parameter is the probability of the target variable being 1 and is determined by the linear combination of the input variables  $X$  and the parameters  $\beta$ . It gives:  $Y_i \sim \text{Bernoulli}(p_i)$ .

By linking function -  $\log(\frac{x}{1-x})$ , the conditional probability of the target variable  $Y$  given the input variables  $X$  is given by <sup>1</sup>:

$$\log\left(\frac{p}{1-p}\right) = X\beta,$$

Where  $p$  is the vector of true predicted probability of  $Y$  being 1,  $X$  is design matrix of the input variables, and  $\beta$  is the vector of the parameters.

$$p = \begin{bmatrix} p_1 \\ p_2 \\ \vdots \\ p_n \end{bmatrix}, \quad X = \begin{bmatrix} X_{11} & X_{12} & \cdots & X_{1p} \\ X_{21} & X_{22} & \cdots & X_{2p} \\ \vdots & \vdots & \ddots & \vdots \\ X_{n1} & X_{n2} & \cdots & X_{np} \end{bmatrix}, \quad \beta = \begin{bmatrix} \beta_1 \\ \beta_2 \\ \vdots \\ \beta_p \end{bmatrix}.$$

After taking the inverse of the link function, we can get the predicted probability of the target variable being 1 <sup>2</sup>:

---

<sup>1</sup>  $\frac{p}{1-p}$  is doing broadcast operation, which makes sure the result is still a vector not linear algebra dot product.

<sup>2</sup>  $\sigma$  is a sigmoid function, which takes  $X$  and  $\beta$  as input and returns the predicted probability, we will use it later for simplicity.

$$p = \sigma(X, \beta) = \frac{1}{1 + e^{-X\beta}}.$$

Based on the above probability model, we can get the likelihood function of parameters  $\beta$  given the data set  $(X, Y)$ :

$$L(\beta|X, Y) = \prod_{i=1}^n p_i^{Y_i} (1 - p_i)^{1-Y_i} \quad (1-1)$$

Taking logarithm of the likelihood function and replacing  $p_i$  with  $\sigma(X_i, \beta)$ <sup>3</sup>, we can get the log-likelihood function of the parameters  $\beta$  given the data set  $(X, Y)$ :

$$\begin{aligned} \ell_n(\beta) &= \sum_{i=1}^n [y_i \log(\sigma(x_i, \beta)) + (1 - y_i) \log(1 - \sigma(x_i, \beta))] \\ &= \sum_{i=1}^n (x_i y_i \beta - \log(1 + e^{x_i \beta})) \end{aligned}$$

We search for the maximum of the log-likelihood function to get the MLE of the parameters  $\beta$ :

$$\hat{\beta} = \arg \max_{\beta} \ell_n(\beta)$$

To find a good solution for the MLE, we can use the Newton-Raphson method to iteratively update the parameters  $\beta$ :

$$\beta_{k+1} = \beta_k - H^{-1}(\beta_k) \nabla \ell_n(\beta_k)$$

Where  $H(\beta_k)$  is the Hessian matrix of the log-likelihood function  $\ell_n(\beta)$  at  $\beta_k$ , and  $\nabla \ell_n(\beta_k)$  is the gradient vector of the log-likelihood function  $\ell_n(\beta)$  at  $\beta_k$ . The initial value of  $\beta$  can be set to any proper value, but a good choice can accelerate the convergence speed.

The Hessian matrix and the gradient vector can be calculated as follows <sup>4</sup>:

---

<sup>3</sup> $X_i$  is the input vector of  $i$ th observation.

<sup>4</sup>We assume the  $X$ 's are independent that we can add  $n$  individual observed Fisher information to get the total Fisher information.  $H$  is summation of  $n$  matrices in  $(p, p)$ ,  $\nabla \ell_n(\beta)$  is the vector of the first derivative of the log-likelihood function

$$H(\beta) = - \sum_{i=1}^n \sigma(x_i, \beta)(1 - \sigma(x_i, \beta))x_i x_i^T$$

$$\nabla \ell_n(\beta) = \sum_{i=1}^n (y_i - \sigma(x_i, \beta))x_i$$

### 1.3 Prior about $\beta$ from bootstrapping method

In situations where prior information about the regression coefficients  $\beta$  is unavailable or limited, a data-driven approach can be employed by using the bootstrap method to approximate the sampling distribution of  $\beta$ . Specifically, we perform repeated sampling with replacement from the observed dataset and estimate  $\beta$  for each bootstrap replicate using a frequentist method such as ordinary least squares. This yields a collection of  $\beta$  estimates, from which we can compute the empirical mean  $\hat{\beta}$  and covariance matrix  $\Sigma_{\text{boot}}$ . These estimates can then be used to define an informative multivariate normal prior for Bayesian inference:

$$\beta \sim \mathcal{N}(\hat{\beta}, \Sigma_{\text{boot}})$$

This bootstrap-based prior reflects the variability observed in the data and provides a pragmatic empirical Bayes approach. While it does not represent a fully Bayesian treatment (as the prior is derived from the data), it enables regularization and uncertainty quantification in a principled way.

### 1.4 Bayesian improved logistic regression

After getting the likelihood function or  $\beta$  from original data and prior distribution of  $\beta$  from bootstrapping method, we can use the Bayes' theorem to get the posterior distribution of  $\beta$ :

$$\mathbb{P}(\beta|X, Y) = \frac{L(\beta|X, Y)p(\beta)}{p(X, Y)} \quad (1-2)$$

Where  $p(X, Y)$  is the marginal likelihood of the data set  $(X, Y)$ .

The posterior does not belong to any known distribution, but it is proportional to its numerator  $p_{\beta_{\text{post}}} = \mathbb{P}(\beta \mid \mathbf{y}, \mathbf{X})$ :

$$\mathbb{P}(\beta \mid \mathbf{y}, \mathbf{X}) \propto \ell_n(\beta) \cdot p(\beta)$$

it gives that we can get the 95% credible interval of  $\beta$  from this proportional unnormalized posterior distribution.

Based on the unnormalized distribution of  $\beta$ , we can get the posterior predictive distribution of the probability of target variable  $Y$  given the input variables  $X$ :

$$\mathbb{P}_{Y_{[\alpha, 1-\alpha]}} = [X\beta_\alpha, X\beta_{1-\alpha}]$$

We set the probability of  $Y$  being 1 as the mean of the posterior predictive distribution of  $Y$ :

$$\mathbb{P}_{Y=1} = E[\beta_{post}X] = \frac{1}{B} \sum_{b=1}^B \sigma(X, \beta_b)$$

This modular design enables robust inference, handles nonlinearities via splines, and leverages empirical Bayesian estimation through bootstrapped priors.

## 1.5 One-VS-Rest classification

We use one-vs-rest method to transit from binary logistic regression to multi-class classification.

The logistic regression model for class  $k$  predicts the probability of an observation belonging to class  $k$  as:

$$P(y = k \mid \mathbf{x}) = \sigma(\mathbf{x}\boldsymbol{\beta}^{(k)}) = \frac{1}{1 + e^{-\mathbf{x}\boldsymbol{\beta}^{(k)}}},$$

where  $\boldsymbol{\beta}^{(k)}$  is the parameter vector for class  $k$ .

After training all  $K$  models, the predicted class for a new observation  $\mathbf{x}$  is determined by selecting the class with the highest predicted probability:

$$\hat{y} = \arg \max_{k \in \{1, 2, \dots, K\}} P(y = k \mid \mathbf{x}).$$

This approach allows binary logistic regression to handle multi-class problems by leveraging the simplicity of binary classification while maintaining interpretability and flexibility.

## 1.6 Implementation of deriving Prior Distribution of $\beta$

To make the modeling work simpler, we only use the data set processed by SMOTE method for work in this section.

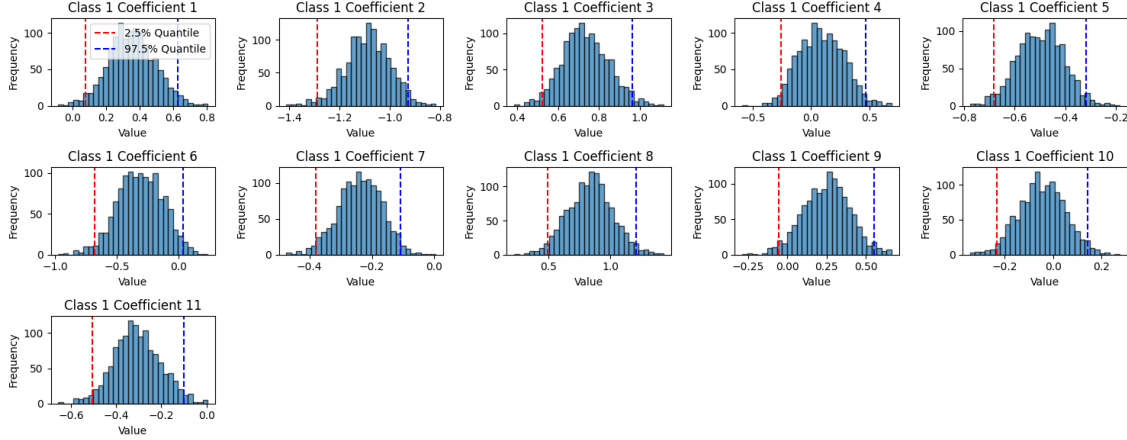


Figure 1: Bootstrap estimates histograms for prior distribution of  $\beta_i$

Figure 1<sup>5</sup> shows the empirical prior distributions of the logistic regression coefficients  $\beta_i$  estimated using bootstrap resampling for class label 1 among the five target categories. Each subplot corresponds to a coefficient in the model, including the intercept term.

The histograms represent the frequency distribution of coefficient estimates across bootstrap samples, giving an approximation of the sampling distribution under repeated sampling. Vertical dashed lines mark the 2.5% and 97.5% quantiles, providing a 95% empirical confidence interval for each  $\beta_i$ . These intervals reflect the uncertainty of the coefficients prior to incorporating the likelihood in the Bayesian inference step.

Figure 2 shows the histogram of values proportional to the posterior distribution of the logistic regression coefficients  $\beta$ . These values represent the product of the likelihood and the bootstrap-estimated prior, up to a normalization constant.

The vertical dashed lines indicate the 10% and 100% quantile thresholds, highlighting the top 90% of posterior-proportional values. This range identifies the most probable  $\beta$  candidates, useful for summarizing the posterior or computing the expected  $\beta$ <sup>6</sup>.

## 1.7 Model performance evaluation

After repeatedly training the model on 5 different classes, we can get the model performance metrics of the models.

<sup>5</sup>Here we only use class 1 in response variable to demonstrate how we obtain the distribution of  $\beta_i$  from bootstrap method, that's there is 'Class 1' string in titles.

<sup>6</sup>The computation is also based on class 1 among the 5 classes.

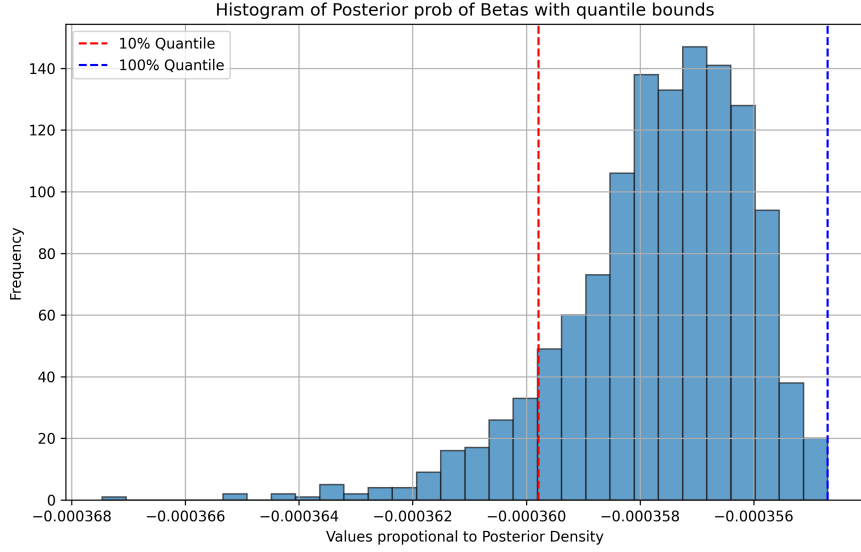


Figure 2: Hist of values proportional to the posterior distribution of  $\beta$  vector

Table 1 shows the model performance metrics of the performance of bayesian logistic regression and MLE logistic regression model on training data.

Table 1: Model evaluation metrics grouped by target

Target	Accuracy	Precision	Recall	F1 Score	Model
1	0.8629	0.6111	0.8800	0.7213	MLE
	0.8629	0.6111	0.8800	0.7213	Bayesian
2	0.6774	0.3729	0.8800	0.5238	MLE
	0.7016	0.3889	0.8400	0.5316	Bayesian
3	0.6694	0.3095	0.5200	0.3881	MLE
	0.6855	0.3158	0.4800	0.3810	Bayesian
4	0.6048	0.2549	0.5417	0.3467	MLE
	0.6774	0.3095	0.5417	0.3939	Bayesian
5	0.7097	0.4000	0.8800	0.5500	MLE
	0.7419	0.4186	0.7200	0.5294	Bayesian

Table 1 compares the performance of MLE and Bayesian logistic regression models on the training dataset across five target classes using standard classification metrics.

Overall, both models exhibit consistent performance on class 1, achieving high accuracy (0.8629), precision (0.6111), recall (0.8800), and F1 score (0.7213), suggesting this class is well modeled.

For the remaining classes, Bayesian models tend to slightly outperform MLE models in terms of accuracy, precision, and F1 score. For example, in class 2, the Bayesian model achieves higher accuracy (0.7016 vs. 0.6774) and better F1 score (0.5316 vs. 0.5238), despite both models showing strong recall (above 0.84). This trend continues in classes 3 and 4, where the Bayesian model improves on accuracy and precision, although the recall values remain close. Class 4 shows a notable improvement in accuracy (0.6774 vs. 0.6048) and F1 score (0.3939 vs. 0.3467) for the Bayesian model.

In class 5, the Bayesian model again achieves higher accuracy (0.7419 vs. 0.7097), although its recall drops slightly compared to MLE (0.7200 vs. 0.8800), leading to a marginally lower F1 score. Overall, the Bayesian approach generally improves model robustness across imbalanced classes by balancing precision and recall more effectively, as reflected in consistently stronger F1 scores in most classes.

*Table 2: Model evaluation metrics on test data*

Target	Accuracy	Precision	Recall	F1 Score	Model
1	0.5887	0.1389	0.2000	0.1639	MLE
	0.5887	0.1389	0.2000	0.1639	Bayesian
2	0.4839	0.1695	0.4000	0.2381	MLE
	0.5081	0.1897	0.4400	0.2651	Bayesian
3	0.5403	0.1190	0.2000	0.1493	MLE
	0.5645	0.1081	0.1600	0.1290	Bayesian
4	0.5565	0.1961	0.4167	0.2667	MLE
	0.5968	0.1905	0.3333	0.2424	Bayesian
5	0.4516	0.1091	0.2400	0.1500	MLE
	0.5323	0.1333	0.2400	0.1714	Bayesian

The evaluation on test data reveals a noticeable drop in performance for both models across all metrics. Accuracy ranges from approximately 45% to 59%, and F1 scores remain low, mostly under 0.27. This indicates challenges in generalizing beyond the training data. Despite the overall low performance, the Bayesian model generally achieves slightly higher F1 scores than MLE, particularly for targets 2, 4, and 5, suggesting better balance between precision and recall under uncertainty.






Figure 3: The importance of each feature in the Bayesian logistic regression model

Figure 3 illustrates the coefficients from five one-vs-rest Bayesian logistic regression models, each targeting a different class. Features with large absolute coefficient values have a stronger influence on the model’s predictions. Notably, features indexed at 0, 1, 5, and 7 exhibit substantial variability across classes, indicating their high discriminative power. For instance, feature 6 is particularly important for Class 1 and Class 4, while feature 5 plays a dominant role in Class 3. Conversely, features 8–10 show relatively small coefficients across all classes, suggesting limited influence on prediction. The direction (sign) of each coefficient also reveals whether a feature increases or decreases the likelihood of the corresponding class.

Figure 4 displays the ROC curves for the five one-vs-rest Bayesian logistic regression models, each classifying one of the five target classes. The Area Under the Curve (AUC) serves as a summary of model performance, where higher values indicate better discrimination. Class 1 achieves the highest AUC (0.94), suggesting strong predictive performance. Class 5 and Class 2 also show good separation with AUCs of 0.85 and 0.82, respectively. In contrast, Class 3 (AUC = 0.70) and especially Class 4 (AUC = 0.67) demonstrate more limited classification ability, with curves closer to the diagonal, indicating greater confusion between classes. Overall, the ROC analysis highlights varying levels of model effectiveness across classes.




Figure 4: ROC curve of each of the Bayesian logistic regression models across 5 classes